

boras-auto-finance-eda

February 23, 2026

0.1 Data Cleaning and Exploratory Data Analysis

0.2 Data Overview :

```
[1]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt
```

```
[2]: data_path = r"D:/Automotive_Loans_dataset/"  
  
files = {  
    "branches": "branch_master.csv",  
    "brokers" : "broker_master.csv",  
    "contracts" : "contracts_master.csv",  
    "dealers" : "dealer_master.csv",  
    "executives" : "executive_master.csv",  
    "repayment" : "repayment_master.csv"  
}  
  
monthly_dpds = {}  
  
for name,file in files.items():  
    monthly_dpds[name] = pd.read_csv(data_path + file)
```

```
branches = monthly_dpds["branches"]  
brokers = monthly_dpds["brokers"]  
contracts = monthly_dpds["contracts"]  
dealers = monthly_dpds["dealers"]  
executives = monthly_dpds["executives"]  
repayment = monthly_dpds["repayment"]
```

```
[22]: print(contracts.head(2),"\n")  
print(contracts.info(),"\n")
```

```
print(contracts.shape, "\n")
```

```
      loan_id customer_name      asset_name invoice_value loan_amount \
0  1000000001     Adira Loke  Mahindra Scorpio        1060000    870000
1  1000000002     Jivika Som  Mahindra Scorpio        1650000   1320000

      emi_amount tenure_months    irr loan_issue_date loan_status \
0       20123           60  13.73    01-05-2023         L
1       46496           36  16.14    18-11-2023         L

      reg_payment_mode psl_tag rc_verified dealer_id broker_id executive_id \
0            ECS        N          N        NaN      NaN    27030005
1            ECS        N          Y        NaN      NaN    27030006

      branch_id manager_id
0       1006    27020412
1       1007    27020412

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9436 entries, 0 to 9435
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_id          9436 non-null   int64  
 1   customer_name    9152 non-null   object  
 2   asset_name       9436 non-null   object  
 3   invoice_value    9436 non-null   int64  
 4   loan_amount      9436 non-null   int64  
 5   emi_amount       9436 non-null   int64  
 6   tenure_months    9436 non-null   int64  
 7   irr              9436 non-null   float64 
 8   loan_issue_date  9436 non-null   object  
 9   loan_status      9436 non-null   object  
 10  reg_payment_mode 8511 non-null   object  
 11  psl_tag          8950 non-null   object  
 12  rc_verified      9169 non-null   object  
 13  dealer_id        3327 non-null   float64 
 14  broker_id        2340 non-null   float64 
 15  executive_id     9436 non-null   int64  
 16  branch_id        9436 non-null   int64  
 17  manager_id       9436 non-null   int64  

dtypes: float64(3), int64(8), object(7)
memory usage: 1.3+ MB
None

(9436, 18)
```

```
[3]: print(repayment.head(2), "\n")
print(repayment.info(), "\n")
print(repayment.shape, "\n")
```

```
    loan_id      due_date    emi_amount  interest_component \
0  1000000001  2023-05-05      20123.0           9954.25
1  1000000001  2023-06-05      20123.0           9837.90

  principal_component    due_amount  amount_received  overdue_amount  dpd_days \
0             10168.75      20123.0          20123.0            0.0         0
1             10285.10      20123.0          20123.0            0.0         0

  balance_amount  branch_id  manager_id
0     859831.25       1006   27020412
1     849546.15       1006   27020412

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 113232 entries, 0 to 113231
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   loan_id          113232 non-null   int64  
 1   due_date         113232 non-null   object 
 2   emi_amount       113232 non-null   float64
 3   interest_component  113232 non-null   float64
 4   principal_component  113232 non-null   float64
 5   due_amount        113232 non-null   float64
 6   amount_received   113232 non-null   float64
 7   overdue_amount    113232 non-null   float64
 8   dpd_days          113232 non-null   int64  
 9   balance_amount    113232 non-null   float64
 10  branch_id         113232 non-null   int64  
 11  manager_id        113232 non-null   int64  
dtypes: float64(7), int64(4), object(1)
memory usage: 10.4+ MB
None

(113232, 12)
```

```
[4]: print(branches.head(2), "\n")
print(branches.info(), "\n")
print(branches.shape, "\n")
```

```
  branch_id  branch_name      state executive_id  executive_name \
0        1001  Mumbai Branch  Maharashtra      27030000  Seher Kulkarni
1        1002  Pune Branch  Maharashtra      27030001  Armaan Bhatia
```

```

manager_id      manager_name
0    27020412  Dharmajan Dhillon
1    27020412  Dharmajan Dhillon

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74 entries, 0 to 73
Data columns (total 7 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   branch_id        74 non-null     int64  
 1   branch_name      74 non-null     object  
 2   state             74 non-null     object  
 3   executive_id     74 non-null     int64  
 4   executive_name   74 non-null     object  
 5   manager_id       74 non-null     int64  
 6   manager_name     74 non-null     object  
dtypes: int64(3), object(4)
memory usage: 4.2+ KB
None

(74, 7)

```

```
[24]: print(brokers.head(), "\n")
print(brokers.info(), "\n")
print(brokers.shape, "\n")
```

	broker_id	broker_name	manager_id	manager_name
0	600000	Aarav Sama	27020209	Badal Kala
1	600001	Ryan Gala	27020772	Vanya Sinha
2	600002	Darshit Suresh	27020666	Vanya Sinha
3	600003	Raunak Kakar	27020666	Anvi Dora
4	600004	Alia Suri	27020772	Dharmajan Dhillon

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   broker_id        200 non-null     int64  
 1   broker_name      200 non-null     object  
 2   manager_id       200 non-null     int64  
 3   manager_name     200 non-null     object  
dtypes: int64(2), object(2)
memory usage: 6.4+ KB
None

```

(200, 4)

```
[5]: print(dealers.head(), "\n")
print(dealers.info(), "\n")
print(dealers.shape, "\n")
```

```
      dealer_id          dealer_name  manager_id  manager_name
0      500000          Bora Ltd Auto    27020562  Uthkarsh Roy
1      500001  Ghose, Dash and Brahmbhatt Motors    27020415  Drishya Raja
2      500002          Vohra Ltd Cars    27020415  Uthkarsh Roy
3      500003  Srinivas, Shetty and Chaudhary Cars    27020047  Uthkarsh Roy
4      500004          Madan LLC Cars    27020412  Purab Sachar
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300 entries, 0 to 299
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   dealer_id        300 non-null    int64  
 1   dealer_name      300 non-null    object  
 2   manager_id       300 non-null    int64  
 3   manager_name     300 non-null    object  
dtypes: int64(2), object(2)
memory usage: 9.5+ KB
None
```

(300, 4)

```
[25]: print(executives.head(), "\n")
print(executives.info(), "\n")
print(executives.shape)
```

```
      executive_id  executive_name  branch_id  manager_id
0      27030000  Seher Kulkarni      1001    27020412
1      27030001  Armaan Bhatia      1002    27020412
2      27030002  Umang Sridhar      1003    27020412
3      27030003  Anya Raval        1004    27020412
4      27030004  Dhanush Dey        1005    27020412
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74 entries, 0 to 73
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
---  --  
 0   executive_id    74 non-null    int64  
 1   executive_name   74 non-null    object  
 2   branch_id       74 non-null    int64
```

```
3    manager_id      74 non-null     int64
dtypes: int64(3), object(1)
memory usage: 2.4+ KB
None

(74, 4)
```

Key Understanding:

- 1. There are 9436 contracts records in Contracts table.
 - 2. There are 113232 records in Repayments table.
 - 3. There are 300 records in Dealer table.
 - 4. There are 200 records in Broker table.
 - 5. There are 74 records in Branch table.
 - 6. There are 74 records in Executive table.
 - 7. It is evident from the column wise value counts that most missing values are present in the Contracts table.
 - 8. We understand that Contracts table has complete customer details and repayment table contains the loan wise repayment records.
-

0.2.1 Finding & handling missing values.

```
[ ]: ## Missing values Assessment.

def summarize_nulls(df,name):
    """
    The given function calculates total null values for each column
    of a table and provides us the summary only for columns that contain
    null values.
    """

    print(f'{-*55}\nSummarized Null Value report for : {name} Table.
    {-*55}')
    null_count = df.isnull().sum()
    null_percent = (null_count/len(df))*100

    summary = pd.DataFrame(
        {
            'column_name' : null_count.index,
            'null_count': null_count.values,
```

```

        'null_percent': null_percent.round(2).astype(str)+"%",
    }
)

if (summary['null_count'].sum() > 0) :

    return summary[summary['null_count']>0].reset_index( drop = True )

else:

    print(f"No Null Values found in {name} Table.")

```

[259]: `print(summarize_nulls(contracts, "Contracts"))
print(summarize_nulls(repayment, "Repayment"))`

Summarized Null Value report for : Contracts Table.

	column_name	null_count	null_percent
0	customer_name	284	3.01%
1	reg_payment_mode	925	9.8%
2	psl_tag	486	5.15%
3	rc_verified	267	2.83%
4	dealer_id	6109	64.74%
5	broker_id	7096	75.2%

Summarized Null Value report for : Repayment Table.

No Null Values found in Repayment Table.

None

[8]: `print(summarize_nulls(branches, "Branches"))
print(summarize_nulls(brokers, "Brokers"))
print(summarize_nulls(dealers, "Dealers"))
print(summarize_nulls(executives, "Executives"))`

Summarized Null Value report for : Branches Table.

No Null Values found in Branches Table.

None

Summarized Null Value report for : Brokers Table.

No Null Values found in Brokers Table.

None

```
Summarized Null Value report for : Dealers Table.
```

```
-----  
No Null Values found in Dealers Table.
```

```
None
```

```
-----  
Summarized Null Value report for : Executives Table.
```

```
-----  
No Null Values found in Executives Table.
```

```
None
```

0.2.2 Key Understanding:

- There are **284 customer names missing out of 9436 contract records but their loan_ids are present.**
- Contract records with missing customer names were retained to preserve repayment history and avoid loss of risk signals.
- For the missing Registered Payment Mode values, PSL tag values, RC verification values, we will have to replace it with “Unknown” as they are business’ regulatory categorizations, hence Mode imputation is not suitable.
- **The null values in Dealer and Broker ID columns will help us in Feature Engineering to create the channel type category.**
- **Logic:** Rows with only dealer_id value will be labelled as “Dealer”, rows with only broker_id value will be labelled as “Broker” and rows with Null values in both will be labelled as “New”.

0.2.3 Data Integrity Validation

```
[ ]: # Checking duplicate loan or payment records, and how should they be handled?
```

```
duplicates_contracts = contracts[contracts.  
    ↪duplicated(subset=['loan_id'],keep=False)]  
print(duplicates_contracts, "\n")  
  
duplicate_repayment = repayment[repayment.  
    ↪duplicated(subset=['loan_id','due_date'],keep=False)]  
print(duplicate_repayment, "\n")  
  
## No duplicate contract records found, tables conform to expected grain.
```

Empty DataFrame

Columns: [loan_id, customer_name, asset_name, invoice_value, loan_amount, emi_amount, tenure_months, irr, loan_issue_date, loan_status, reg_payment_mode, psl_tag, rc_verified, dealer_id, broker_id, executive_id, branch_id, manager_id]

```
Index: []
```

```
Empty DataFrame  
Columns: [loan_id, due_date, emi_amount, interest_component,  
principal_component, due_amount, amount_received, overdue_amount, dpd_days,  
balance_amount, branch_id, manager_id]  
Index: []
```

```
[6]: print(f"Unique Values for loan_id in repayments: {repayment['loan_id'] .  
↳nunique()}\n")  
print(f"Unique Values for loan_id in contracts: {contracts['loan_id'] .  
↳nunique()}\n")  
  
## Repayment records exist for all contracts in both tables. No orphan records  
↳found.
```

```
Unique Values for loan_id in repayments: 9436
```

```
Unique Values for loan_id in contracts: 9436
```

```
[7]: print(f"Unique Values for dealer_id in contracts: {contracts['dealer_id'] .  
↳nunique()}\n")  
print(f"Unique Values for dealer_id in dealrs: {dealers['dealer_id'] .  
↳nunique()}\n")  
print(f"Unique Values for broker_id in contracts: {contracts['broker_id'] .  
↳nunique()}\n")  
print(f"Unique Values for broker_id in brokers: {brokers['broker_id'] .  
↳nunique()}\n")  
  
# All registered dealers/brokers have at least one contract in the portfolio.
```

```
Unique Values for dealer_id in contracts: 300
```

```
Unique Values for dealer_id in dealrs: 300
```

```
Unique Values for broker_id in contracts: 200
```

```
Unique Values for broker_id in brokers: 200
```

0.3 ##### Feature Engineering

```
[31]: contracts.head(2) # Before
```

```
[31]:      loan_id customer_name          asset_name invoice_value loan_amount \
0  1000000001     Adira Loke  Mahindra Scorpio      1060000    870000
1  1000000002     Jivika Som  Mahindra Scorpio      1650000   1320000

      emi_amount tenure_months      irr loan_issue_date loan_status \
0        20123           60  13.73    01-05-2023         L
1        46496           36  16.14    18-11-2023         L

      reg_payment_mode psl_tag rc_verified dealer_id broker_id executive_id \
0            ECS       N            N        NaN       NaN    27030005
1            ECS       N            Y        NaN       NaN    27030006

      branch_id manager_id
0        1006    27020412
1        1007    27020412
```

```
[3]: # Created channel_type column to track contract lead source existing/channel/
    ↵broker for downstream analysis
```

```
contracts['channel_type'] = np.where(
    contracts['dealer_id'].notna() & contracts['broker_id'].isna(), 'Dealer',
    np.where(contracts['broker_id'].notna() & contracts['dealer_id'].isna(), 'Broker', 'Self')
)
```

```
[263]: contracts.head(2) #After
```

```
[263]:      loan_id customer_name          asset_name invoice_value loan_amount \
0  1000000001     Adira Loke  Mahindra Scorpio      1060000    870000
1  1000000002     Jivika Som  Mahindra Scorpio      1650000   1320000

      emi_amount tenure_months      irr loan_issue_date loan_status \
0        20123           60  13.73    01-05-2023         L
1        46496           36  16.14    18-11-2023         L

      reg_payment_mode psl_tag rc_verified dealer_id broker_id executive_id \
0            ECS       N            N        NaN       NaN    27030005
1            ECS       N            Y        NaN       NaN    27030006

      branch_id manager_id channel_type
0        1006    27020412        Self
1        1007    27020412        Self
```

```
[4]: ## Created a 'dpd_bucket' column to classify contracts into Regular/X bucket/S1/  
    ↪S2/S3/NPA contracts for downstream analysis.
```

```
conditions = [  
  
    (repayment['dpd_days'] == 0 ),  
    (repayment['dpd_days'] > 0 ) & (repayment['dpd_days'] < 30 ),  
    (repayment['dpd_days'] >= 30 ) & ( repayment['dpd_days'] < 60 ),  
    (repayment['dpd_days'] >= 60 ) & ( repayment['dpd_days'] < 90 ),  
    (repayment['dpd_days'] >= 90)  
]  
  
classification = ['Regular', 'X bucket', 'S1', 'S2', 'NPA' ]  
  
repayment['dpd_bucket'] = np.select(conditions,classification,default=  
    ↪'Unknown')
```

```
[5]: repayment.head(2)
```

```
[5]:      loan_id   due_date   emi_amount   interest_component  \  
0  1000000001  2023-05-05     20123.0          9954.25  
1  1000000001  2023-06-05     20123.0          9837.90  
  
      principal_component   due_amount   amount_received   overdue_amount   dpd_days  \  
0            10168.75     20123.0           20123.0            0.0          0  
1            10285.10     20123.0           20123.0            0.0          0  
  
      balance_amount   branch_id   manager_id   dpd_bucket  
0        859831.25       1006     27020412   Regular  
1        849546.15       1006     27020412   Regular
```

```
[6]: # Payment Delay Patterns by EMI Sequence  
# Checking if later EMIs show higher delinquency than early the EMIs.
```

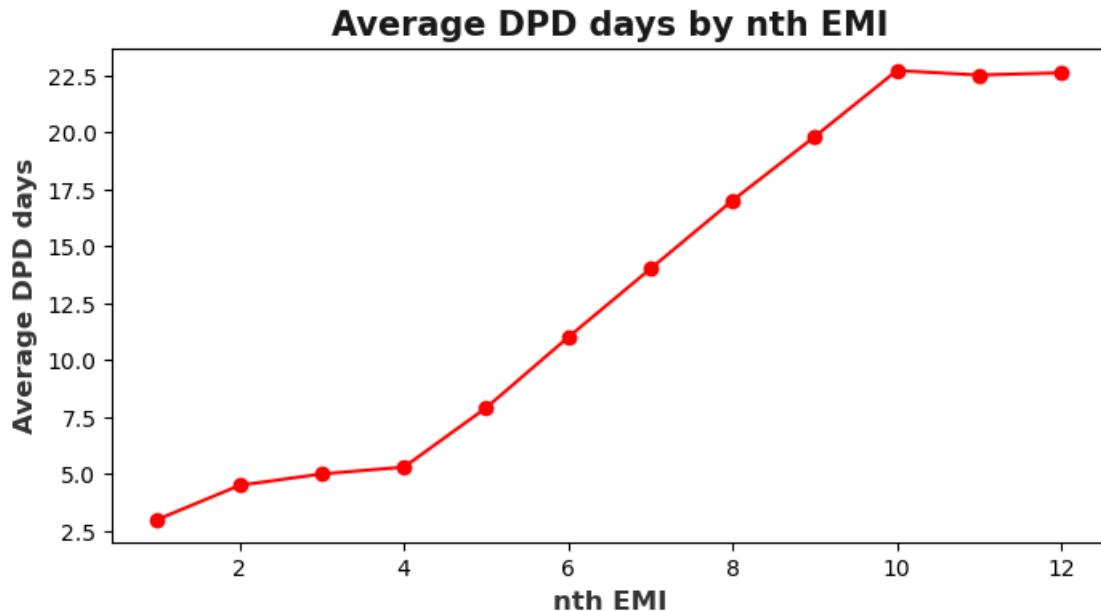
```
max_date = repayment['due_date'].max()  
snapshot = repayment[repayment['due_date'] <= max_date]  
snapshot['emi_sequence'] = snapshot.groupby('loan_id').cumcount()+1  
# Added serial numbers to each of the contract's EMI.  
snapshot.head()  
  
sequence_dpd = snapshot.groupby('emi_sequence')['dpd_days'].  
    ↪agg(['mean','count']).round(1)  
# Averages and count of DPD days for all nth EMIs .  
print("Average DPD days by EMI Sequence:\n")  
print(sequence_dpd)
```

Average DPD days by EMI Sequence:

	mean	count
emi_sequence		
1	3.0	9436
2	4.5	9436
3	5.0	9436
4	5.3	9436
5	7.9	9436
6	11.0	9436
7	14.0	9436
8	17.0	9436
9	19.8	9436
10	22.7	9436
11	22.5	9436
12	22.6	9436

```
[271]: plt.figure(figsize=(8,4))
plt.plot(sequence_dpd.index,sequence_dpd['mean'],marker='o',color = 'red')
plt.title("Average DPD days by nth EMI",fontsize = 15, fontweight = "bold",
          alpha = 0.9)
plt.xlabel("nth EMI",fontsize = 12, fontweight = "bold",alpha = 0.8)
plt.ylabel("Average DPD days",fontsize = 12, fontweight = "bold",alpha = 0.8)
```

```
[271]: Text(0, 0.5, 'Average DPD days')
```



```
[8]: ## Checking if there is correlation between loan tenure and DPD buckets
    ↵(Regular/S1/S2/NPA).

print(f"---*65) \nColumns of repayment table: \n---*65)\n {repayment.
    ↵columns}\n")
print(f"---*65) \nLast date in the repayment table is '{repayment['due_date'].
    ↵max()}'\n---*65)")

repayment['due_date'] = pd.to_datetime(repayment['due_date'])
```

Columns of repayment table:

```
Index(['loan_id', 'due_date', 'emi_amount', 'interest_component',
       'principal_component', 'due_amount', 'amount_received',
       'overdue_amount', 'dpd_days', 'balance_amount', 'branch_id',
       'manager_id', 'dpd_bucket'],
      dtype='object')
```

Last date in the repayment table is '2024-11-15'

```
[9]: november_2024 = repayment[
    (repayment['due_date'].dt.year == 2024) &
    (repayment['due_date'].dt.month == 11)
][['loan_id', 'dpd_bucket']]

print(f"Total contracts as on November 2024: {november_2024.shape}\n")

# We have total 839 loan contracts as on November 2024.

corr_monthly_dpd = contracts[['loan_id', 'tenure_months']].merge(
    november_2024, on = 'loan_id', how = 'inner'
)

corr_monthly_dpd.head()
```

Total contracts as on November 2024: (839, 2)

```
[9]:      loan_id  tenure_months  dpd_bucket
0  1000000007          36    Regular
1  1000000016          60    Regular
2  1000000019          60    Regular
```

```

3 1000000024          36    Regular
4 1000000030          60    Regular

```

[]: # Numeric bucket for correlation.

```

dpd_bucket_map = {'Regular':0,'S1':1,'S2':2,'NPA': 3}

corr_monthly_dpd['dpd_bucket_num'] = corr_monthly_dpd['dpd_bucket'].
    ↪map(dpd_bucket_map)
corr_monthly_dpd.head()

```

[]: loan_id tenure_months dpd_bucket dpd_bucket_num

loan_id	tenure_months	dpd_bucket	dpd_bucket_num
0 1000000007	36	Regular	0
1 1000000016	60	Regular	0
2 1000000019	60	Regular	0
3 1000000024	36	Regular	0
4 1000000030	60	Regular	0

[300]: correlation = corr_monthly_dpd[['tenure_months','dpd_bucket_num']].corr(method ↪= 'spearman')

print(f"{'-'*45}\nCorrelation between Tenure and DPD buckets: \n{'-'*45}")
print(correlation, "\n")

rho = correlation.loc[['tenure_months','dpd_bucket_num']]
print(f"{'-'*45}\nSpearman = {rho:.3f} \n{'-'*45}")

Correlation between Tenure and DPD buckets:

	tenure_months	dpd_bucket_num
tenure_months	1.000000	0.005343
dpd_bucket_num	0.005343	1.000000

Spearman = 0.005

[]: # Monthly Average DPD days trend by year:

```

monthly_dpd = repayment.groupby([
    repayment['due_date'].dt.year.rename('year'),
    repayment['due_date'].dt.month.rename('month')
])['dpd_days'].mean().reset_index(name='avg_dpd')

```

```

monthly_dpd['month_label'] = monthly_dpd['month'].map({

```

```

    1:'Jan', 2:'Feb', 3:'Mar', 4:'Apr', 5:'May', 6:'Jun',
    7:'Jul', 8:'Aug', 9:'Sep', 10:'Oct', 11:'Nov', 12:'Dec'
})

print(monthly_dpd)

```

	year	month	avg_dpd	month_label
0	2023	1	2.951432	Jan
1	2023	2	4.055118	Feb
2	2023	3	4.375267	Mar
3	2023	4	4.259674	Apr
4	2023	5	4.927866	May
5	2023	6	6.230134	Jun
6	2023	7	7.263330	Jul
7	2023	8	8.428300	Aug
8	2023	9	9.497890	Sep
9	2023	10	10.957650	Oct
10	2023	11	12.440386	Nov
11	2023	12	13.210047	Dec
12	2024	1	14.362331	Jan
13	2024	2	15.686299	Feb
14	2024	3	16.279791	Mar
15	2024	4	17.532097	Apr
16	2024	5	18.246126	May
17	2024	6	19.792241	Jun
18	2024	7	19.604793	Jul
19	2024	8	20.051315	Aug
20	2024	9	21.410146	Sep
21	2024	10	21.608392	Oct
22	2024	11	20.059595	Nov

```
[186]: for_2023 = monthly_dpd[monthly_dpd['year']==2023][['month','month_label','avg_dpd']]
for_2024 = monthly_dpd[monthly_dpd['year']==2024][['month','month_label','avg_dpd']]
```

```
[219]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(for_2023['month_label'],for_2023['avg_dpd'],marker = 'o', color = 'red')
plt.xlabel('Months',fontweight = 'bold', alpha = 0.8)
plt.ylabel('Average DPD days',fontweight = 'bold', alpha = 0.8)
plt.title('Monthly Average DPD days Trend for year 2023',fontsize = 13,fontweight = 'bold', alpha = 0.9)

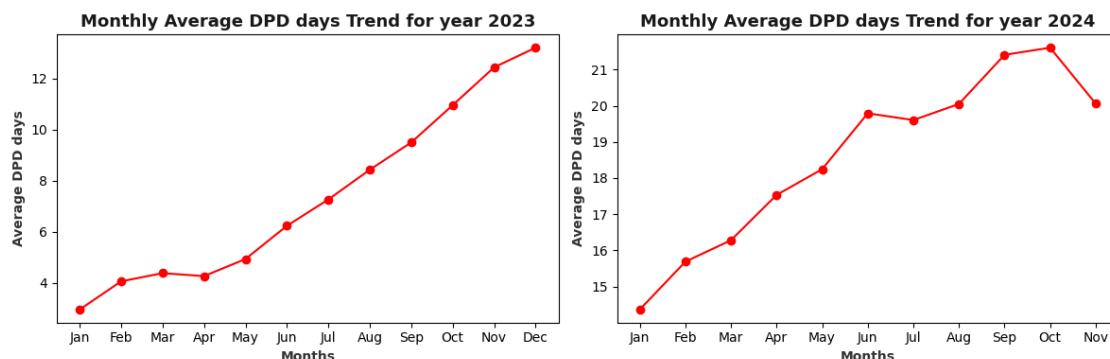
plt.subplot(1,2,2)
```

```

plt.plot(for_2024['month_label'],for_2024['avg_dpd'],marker = 'o', color = 'red')
plt.xlabel('Months',fontweight = 'bold', alpha = 0.8)
plt.ylabel('Average DPD days',fontweight = 'bold', alpha = 0.8)
plt.title('Monthly Average DPD days Trend for year 2024',fontsize = 13,fontweight = 'bold', alpha = 0.9)

plt.tight_layout()
plt.show()

```



[]: ## Loading DataFrames into MySQL.

```

from sqlalchemy import create_engine

engine = create_engine(
    "mysql+pymysql://root:MYsql%4060323@localhost:3306/automotive_loans"
)

```

[18]: tables = {

```

    "branches"      : branches,
    "brokers"       : brokers,
    "contracts"     : contracts,
    "dealers"       : dealers,
    "executives"    : executives,
    "repayment"     : repayment
}

```

```

for table_name, df in tables.items():

    df.to_sql(
        name = table_name,
        con = engine,

```

```
        if_exists = "replace",
        index = False
    )
    print(f"{table_name} loaded.\n")
```

branches loaded.

brokers loaded.

contracts loaded.

dealers loaded.

executives loaded.

repayment loaded.